Model 1 Selection Report

Elias Cheleuitte, Micaleb Johnson, Morgan Williams

Data Cleaning:

As we first began trying to model players' career win shares, we ran into an issue of a number of missing values. All but one of the highschool players and roughly a third of the international players were missing information for turnovers per game prior to being drafted, because of this we decided to remove the turnover variable from any potential models. After removing turnovers from our model selection process we still were unable to properly model our data due to missing data for 9 players spread between three variables: assists per game, steals per game, and blocks per game. Seven of these were high school players, while the other two were international players. Due to our large original sample size of 468 we concluded that removing the 9 players from our data was best. Furthermore, we removed two more players from our data Qyntel Woods and Kedrick Brown because, while they played in college in the United States, they played at Junior Colleges and not the NCAA like the rest of our college players.

Model Selection:

Once our data issues were corrected, we randomly split our original dataset by rows into train and test sets, with 80% of our data being put into the training set and the remaining 20% into the testing set. Now working only with the training dataset, we automated the selection of the optimal regression model based on the AIC and BIC for all possible model subsets starting from a full model with all 15 covariates down to an intercept only model. As a result, a six variable model consisting of Pick, Age, Height (inches), Weight, pre NBA draft field goal percentage (The number of shots a player makes / The number of shots a player attempts), and pre NBA draft steals per game, looked to be the best model based on our metrics. All of these covariates had p-values less than the significance level of .05 and thus all are likely meaningful to the model. Another model with seven variables including the same covariates as the previous six variable model plus pre NBA free throw percentage also looked promising, with all but one of the variables appearing to be significant by their p-values. We also looked for any interaction effects to include but none appeared to drastically improve the models performance without greatly increasing the complexity of the model thus none were added.

Figure 1: Best models sorted by average AIC and BIC rank

	model	Num.var	loglik	aic	bic	Residual.Deviance	Residual.df		Null.Deviance	Null.df	Dev.pval	aic.rank	bic.rank	avg.rank
26041	ws ~ 1 + Pick + Age + Inches + NCAA_fgpct + NCAA_TRB + NCAA_STL + NCAA_BLK	7	-1725.563	3469.126	3504.225	272981.6	357	26041	344106	364	1	22	15	18.5
26033	ws ~ 1 + Pick + Age + Inches + Weight + NCAA_fgpct + NCAA_TRB + NCAA_STL + NCAA_BLK	8	-1723.333	3466.666	3505.665	269666.2	356	26033	344106	364	1	1	37	19.0
30385	ws ~ 1 + Pick + Age + Inches + Weight + NCAA_fgpct + NCAA_ftpct + NCAA_STL	7	-1725.860	3469.721	3504.820	273426.7	357	30385	344106	364	1	31	21	26.0
30641	ws ~ 1 + Pick + Age + Inches + Weight + NCAA_fgpct + NCAA_STL	6	-1727.354	3470.709	3501.908	275674.6	358	30641	344106	364	1	66	3	34.5
25021	ws ~ 1 + Pick + Age + NCAA_fgpct + NCAA_TRB + NCAA_AST + NCAA_STL + NCAA_BLK	7	-1726.085	3470.169	3505.268	273763.1	357	25021	344106	364	1	43	30	36.5
26035	ws ~ 1 + Pick + Inches + Weight + NCAA_fgpct + NCAA_TRB + NCAA_STL + NCAA_BLK	7	-1726.149	3470.298	3505.397	273859.9	357	26035	344106	364	1	49	32	40.5
14257	ws ~ 1 + Pick + Age + Inches + Weight + NCAA_fgpct + NCAA_STL + NCAA_games	7	-1726.377	3470.753	3505.852	274201.4	357	14257	344106	364	1	68	43	55.5
30513	ws ~ 1 + Pick + Age + Inches + Weight + NCAA_fgpct + NCAA_3ptpct + NCAA_STL	7	-1726.393	3470.787	3505.886	274226.4	357	30513	344106	364	1	71	45	58.0
30129	ws ~ 1 + Pick + Age + Inches + Weight + NCAA_fgpct + NCAA_TRB + NCAA_STL	7	-1726.456	3470.912	3506.011	274320.5	357	30129	344106	364	1	84	48	66.0
26545	ws ~ 1 + Pick + Age + Inches + Weight + NCAA_fgpct + NCAA_STL + NCAA_BLK	7	-1726.458	3470.916	3506.015	274324.0	357	26545	344106	364	1	85	49	67.0
29617	ws ~ 1 + Pick + Age + Inches + Weight + NCAA_fgpct + NCAA_AST + NCAA_STL	7	-1726.539	3471.078	3506.178	274445.9	357	29617	344106	364	1	94	51	72.5
25017	ws ~ 1 + Pick + Age + Inches + NCAA_fgpct + NCAA_TRB + NCAA_AST + NCAA_STL + NCAA_BLK	8	-1724.822	3469.645	3508.644	271876.0	356	25017	344106	364	1	29	120	74.5
14001	ws ~ 1 + Pick + Age + Inches + Weight + NCAA_fgpct + NCAA_ftpct + NCAA_STL + NCAA_games	8	-1724.913	3469.827	3508.826	272011.8	356	14001	344106	364	1	35	129	82.0
9657	ws ~ 1 + Pick + Age + Inches + NCAA_fgpct + NCAA_TRB + NCAA_STL + NCAA_BLK + NCAA_games	8	-1724.941	3469.881	3508.880	272052.3	356	9657	344106	364	1	38	134	86.0
29873	ws ~ 1 + Pick + Age + Inches + Weight + NCAA_fgpct + NCAA_ftpct + NCAA_TRB + NCAA_STL	8	-1725.003	3470.006	3509.005	272145.2	356	29873	344106	364	1	41	143	92.0
26043	ws ~ 1 + Pick + Inches + NCAA_fgpct + NCAA_TRB + NCAA_STL + NCAA_BLK	6	-1728.087	3472.174	3503.373	276782.9	358	26043	344106	364	1	183	9	96.0
25009	ws ~ 1 + Pick + Age + Inches + Weight + NCAA_fgpct + NCAA_TRB + NCAA_AST + NCAA_STL + NCAA_BLK	9	-1722.465	3466.930	3509.829	268386.9	355	25009	344106	364	1	2	195	98.5
29361	ws ~ 1 + Pick + Age + Inches + Weight + NCAA_fgpct + NCAA_ftpct + NCAA_AST + NCAA_STL	8	-1725.116	3470.231	3509.230	272313.4	356	29361	344106	364	1	46	157	101.5
8637	ws ~ 1 + Pick + Age + NCAA_fgpct + NCAA_TRB + NCAA_AST + NCAA_STL + NCAA_BLK + NCAA_games	8	-1725.156	3470.312	3509.311	272373.8	356	8637	344106	364	1	50	161	105.5
9649	ws ~ 1 + Pick + Age + Inches + Weight + NCAA_fgpct + NCAA_TRB + NCAA_STL + NCAA_BLK + NCAA_games	9	-1722.592	3467.184	3510.083	268573.6	355	9649	344106	364	1	3	215	109.0
25777	ws ~ 1 + Pick + Age + Inches + Weight + NCAA_fgpct + NCAA_ftpct + NCAA_TRB + NCAA_STL + NCAA_BLK	9	-1722.643	3467.286	3510.185	268649.0	355	25777	344106	364	1	4	222	113.0
25785	ws ~ 1 + Pick + Age + Inches + NCAA_fgpct + NCAA_ftpct + NCAA_TRB + NCAA_STL + NCAA_BLK	8	-1725.229	3470.458	3509.457	272482.4	356	25785	344106	364	1	58	168	113.0

Before going forward in determining which of the two models would be better for predicting NBA career win shares, we checked the necessary plots of both models to insure that they met the assumptions of a linear regression model. As you can see below the residual plots for both models do not appear to have

any sort of pattern, neither have any highly influential points, nor are either drastically non normal and thus we concluded that both models met the assumptions.

Figure 2: Plots of the six variable model

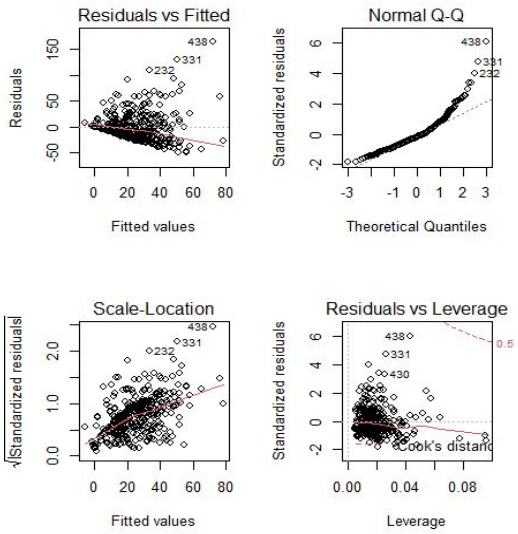

Figure 3: Plots of the seven variable model

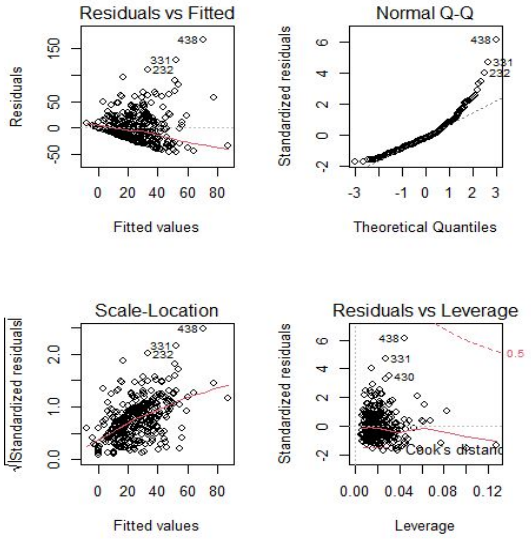

Figure 4: Six variable model

```
Call:
lm(formula = ws ~ 1 + Pick + Age + Inches + Weight + NCAA_fgpct +
    NCAA_STL, data = train)

Residuals:
   Min     1Q Median     3Q    Max
-50.15 -17.42  -5.39  10.86 164.30

Coefficients:
            Estimate Std. Error t value Pr(>|t|)
(Intercept) 196.0927    51.0653   3.840 0.000145 ***
Pick         -0.8997     0.1830  -4.916 1.35e-06 ***
Age          -2.6013     1.1655  -2.232 0.026233 *
Inches       -2.7691     0.7470  -3.707 0.000243 ***
Weight        0.2795     0.1028   2.718 0.006893 **
NCAA_fgpct   86.9358    29.7422   2.923 0.003687 **
NCAA_STL      8.9628     2.3232   3.858 0.000136 ***
---
Signif. codes:  0 ‘***’ 0.001 ‘**’ 0.01 ‘*’ 0.05 ‘.’ 0.1 ‘ ’ 1

Residual standard error: 27.75 on 358 degrees of freedom
Multiple R-squared:  0.1989,    Adjusted R-squared:  0.1854
F-statistic: 14.81 on 6 and 358 DF,  p-value: 3.905e-15
```

Figure 5: Seven variable model

```
Call:
lm(formula = ws ~ 1 + Pick + Age + Inches + Weight + NCAA_fgpct +
    NCAA_ftpct + NCAA_STL, data = train)

Residuals:
    Min      1Q  Median      3Q     Max
-46.728 -17.085  -5.275  11.589 166.101

Coefficients:
            Estimate Std. Error t value Pr(>|t|)
(Intercept) 164.5176    54.1604   3.038 0.002560 **
Pick         -0.8583     0.1841  -4.662 4.44e-06 ***
Age          -2.8831     1.1739  -2.456 0.014525 *
Inches       -2.7156     0.7456  -3.642 0.000311 ***
Weight        0.3039     0.1035   2.935 0.003551 **
NCAA_fgpct   96.8072    30.2166   3.204 0.001478 **
NCAA_ftpct   31.6313    18.4638   1.713 0.087553 .
NCAA_STL      8.6381     2.3247   3.716 0.000235 ***
---
Signif. codes:  0 ‘***’ 0.001 ‘**’ 0.01 ‘*’ 0.05 ‘.’ 0.1 ‘ ’ 1

Residual standard error: 27.67 on 357 degrees of freedom
Multiple R-squared:  0.2054,    Adjusted R-squared:  0.1898
F-statistic: 13.18 on 7 and 357 DF,  p-value: 3.978e-15
```

When we looked at the adjusted r^2 values for both models we saw that there was little difference between them. The six variable model had an adjusted r^2 of .1854 whereas the seven variable model's was only slightly higher at .1898, while these values seem low it is to be expected when analyzing a problem with as many complicating factors as we are. We wanted to verify that there was not a significant difference between the two models. So we performed an ANOVA test between the two models on our training data. Our ANOVA test between the two models provided p value 0.08755. Based on this result we failed to reject Ho, so we concluded that the more complex model wasn't a significantly better fit on the train data than the less complex six variable model, thus we should use the six variable model.

Figure 6: Anova test on six and seven variable models

```
Analysis of Variance Table

Model 1: ws ~ 1 + Pick + Age + Inches + Weight + NCAA_fgpct + NCAA_STL
Model 2: ws ~ 1 + Pick + Age + Inches + Weight + NCAA_fgpct + NCAA_ftpct +
    NCAA_STL
  Res.Df    RSS Df Sum of Sq      F  Pr(>F)
1    358 275675
2    357 273427  1    2247.8 2.9349 0.08755 .
---
Signif. codes:  0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1
```

Model Performance:

In order to correctly assess how effective our model is in general, we needed to test its performance on unseen data. Thus, we applied the models to the test dataset that was set aside before our model selection process began.

Figure 7: Plots of six variable model on test data

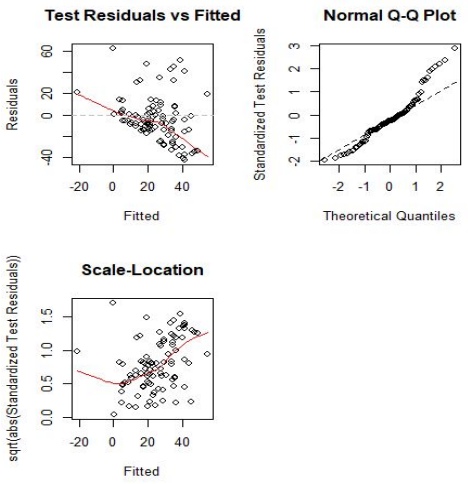

Figure 8: Plots of seven variable model on test data

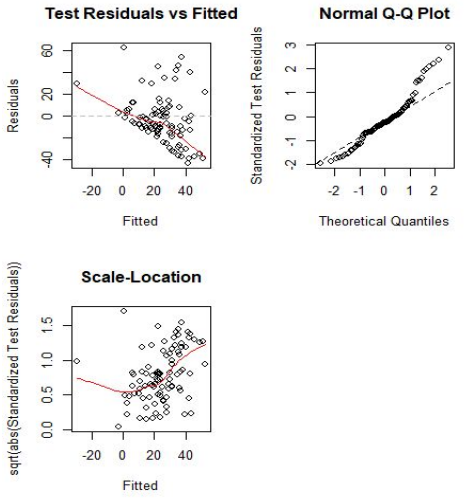

The residual plots for the six and seven variable models with our test data are very similar and only vary slightly. We need to look further into why these plots are not completely satisfying the assumptions for regression, as the residual plots appear to have some pattern and the normality quantile plot splits from the line in the upper quantile. In order to evaluate which model fit the test data better we used Mean Squared Prediction Error. For our six variable model we got a mean squared prediction error of 471.6991 and for the seven variable model we got a mean squared prediction error of 476.1226 and since a lower value indicates a better fitting model, the six variable model is better.

Overall the six variable model appears to be nearly as accurate as the seven variable model without the added complexity of an additional variable and it is the better model for projecting the career win shares of NBA players.

References:

Leaps, RSQLite, kader, rvest, tidyverse, knitr, goftest, ggplot2, ggiraphExtra, glmulti